

**BITS** Pilani  
Pilani Campus

# Network Programming

K Hari Babu  
Department of Computer Science & Information Systems



# Outline

# Outline

- Advanced I/O
  - `recv()`, `send()`
  - `readv()`, `writev()`
  - `recvmsg()`, `sendmsg()`



# Advanced I/O Functions

(T1: ch 14)

# recv() and send()

- The recv() and send() system calls perform I/O on connected sockets (TCP or connected UDP sockets).
- Socket-Specific I/O System Calls:
  - They provide socket-specific functionality not available with read(0) and write().

```
1 #include <sys/socket.h>
2 ssize_t recv(int sockfd , void * buffer , size_t length , int flags );
3 //Returns number of bytes received, 0 on EOF, or -1 on error
4 ssize_t send(int sockfd , const void * buffer , size_t length , int flags );
5 //Returns number of bytes sent, or -1 on error
```

- Same as read() and write() except for *flags*.
- Return values are same as read() and write().

# recv() flags

- **MSG\_DONTWAIT:**
  - perform a non-blocking recv().
  - can be done using fcntl() call but that will make sock fd non-blocking. Here only this operation is non-blocking.
- **MSG\_OOB:**
  - receive out-of-band data on the socket.
- **MSG\_PEEK:**
  - retrieve a copy of the requested bytes from the socket buffer.
  - Data is not removed from the socket buffer.
  - Used for knowing the no of bytes available on the buffer.
- **MSG\_WAITALL**
  - Blocks until *length* bytes are read from socket buffer.
  - May get interrupted by signals.

# send() flags


- **MSG\_DONTWAIT**
  - Perform a non-blocking send.
- **MSG\_MORE**
  - Data written using send() or sendto() calls with this flag is packaged into a single datagram until a send() without this flag.
- **MSG\_NOSIGNAL**
  - Do not generate SIGPIPE signal. Return only EPIPE error.
- **MSG\_OOB**
  - Write out of band data on TCP.

# sendfile() sys call (R1: 61.4)


- Transferring large file in web servers requires repeated calls to read() and write().
  - This is inefficient.

```

1 while ((n = read(diskfilefd, buf, BUZ_SIZE)) > 0)
2     write(sockfd, buf, n);
  
```



a) *read() + write()*



b) *sendfile()*

# sendfile() sys call (R1:61.4)



- The sendfile() sys call is designed to eliminate copying file data into user space.
  - File contents are directly transferred to the socket without going through user space.
  - This is referred as a *zero-copy transfer*.

```
1 #include <sys/sendfile.h>
2 ssize_t sendfile(int out_fd, int in_fd, off_t * offset, size_t count );
3 //Returns number of bytes transferred, or -1 on error
```

- out\_fd*: is the socket fd.
- In\_fd*: is regular file fd.
- off\_t*: is the offset. This is a value-result argument.
- count* is the number of bytes to be transferred.
- sendfile* doesn't change the file offset for *in\_fd*.

# readv() and writev()

- The readv() and writev() system calls perform scatter-gather I/O.
- iov* points to an array of buffers, each in *iovec* structure.

```


1 #include <sys/uio.h>
2 ssize_t readv(int fd , const struct iovec * iov , int iovcnt );
3 //Returns number of bytes read, 0 on EOF, or -1 on error
4 ssize_t writev(int fd , const struct iovec * iov , int iovcnt );
5 //Returns number of bytes written, or -1 on error

```

```

8 struct iovec {
9     void *iov_base; /* Start address of buffer */
10    size_t iov_len; /* Number of bytes to transfer to/from buffer */
11 };

```



# readv() and writev()

- The readv() system call performs scatter input:
  - Reads from the file and puts the data into the buffer starting at iov[0]. Once the first buffer is full, it goes to another.
- readv() completes atomically:
  - Kernel performs a single data transfer.
  - Assured that all the bytes read are contiguous in the file. File offset can't be changed by other process.
- The writev() call performs gather output:
  - Starting from the first buffer, writes the data contiguously into the file.
  - Partial write is possible.
- writev() completes atomically.
- readv() and writev() are used for convenience and speed.
  - Reduce number of sys calls.

# sendmsg() & recvmsg() sys calls

- The sendmsg() and recvmsg() system calls are the most general purpose of the socket I/O system calls.
  - The sendmsg() system call can do everything that is done by write(), send(), and sendto();
  - the recvmsg() system call can do everything that is done by read(), recv(), and recvfrom().

```

1 #include <sys/socket.h>
2 ssize_t recvmsg(int sockfd, struct msghdr *msg, int flags);
3 ssize_t sendmsg(int sockfd, struct msghdr *msg, int flags);
4 //Both return: number of bytes read or written if OK, -1 on error

```

```

1 struct msghdr {
2     void          *msg_name;           /* protocol address */
3     socklen_t     msg_namelen;        /* size of protocol address */
4     struct iovec  *msg_iov;           /* scatter/gather array */
5     int           msg iovlen;         /* # elements in msg_iov */
6     void          *msg_control;       /* ancillary data (cmsghdr struct) */
7     socklen_t     msg_controllen;    /* length of ancillary data */
8     int           msg_flags;          /* flags returned by recvmsg() */
9 };

```

# sendmsg() & recvmsg() sys calls

- Can be used to send or receive ancillary data (control information).
- Flags are

Flag	Examined by: Send flags Sendto flags Sendmsg flags	Examined by: recv flags recvfrom flags recvmsg flags	Returned by: Recvmsg msg_flags
MSG_DONTROUTE	•		
MSG_DONTWAIT	•	•	
MSG_PEEK		•	
MSG_WAITALL		•	
MSG_EOR	•		•
MSG_OOB	•	•	•
MSG_BCAST			•
MSG_MCAST			•
MSG_TRUNC			•
MSG_CTRUNC			•


# Flags returned by rcvmsg()

- **MSG\_BCAST**
  - is returned if the datagram was received as a broadcast.
- **MSG\_MCAST**
  - is returned if the datagram was received as a link-layer multicast.
- **MSG\_TRUNC**
  - is returned if the datagram was truncated
- **MSG\_CTRUNC**
  - is returned if the ancillary data was truncated
- **MSG\_EOR**
  - is turned on if the returned data ends a logical record.
- **MSG\_OOB**
  - This flag is never returned for TCP out-of-band data. This flag is returned by other protocol suites (e.g., the OSI protocols).
- **MSG\_NOTIFICATION**
  - This flag is returned for SCTP receivers to indicate that the message read is an event notification, not a data message.

# Ancillary Data



- Ancillary data can be sent and received using the `msg_control` and `msg_controllen` members of the `msghdr` structure.
  - Another term for ancillary data is control information.



```
1 struct cmsghdr {  
2     socklen_t cmsg_len; /* length in bytes, including this structure */  
3     int cmsg_level; /* originating protocol */  
4     int cmsg_type; /* protocol-specific type */  
5     /* followed by unsigned char cmsg_data[] */  
6 };
```


# Ancillary Data

- Ancillary data is domain specific.

<b>Protocol</b>	<b>cmsg_level</b>	<b>Cmsg_type</b>	<b>Description</b>
IPv4	IPPROTO_IP	IP_RECVSTADD R IP_RECVIF	receive destination address with UDP datagram receive interface index with UDP datagram
IPv6	IPPROTO_IPV6	IPV6_DSTOPTS IPV6_HOPLIMIT IPV6_HOPOPTS IPV6_NEXTHOP IPV6_PKTINFO IPV6_RTHDR	specify / receive destination options specify / receive hop limit specify / receive hop-by-hop options specify next-hop address specify / receive packet information specify / receive routing header
Unix domain	SOL_SOCKET	SCM_RIGHTS SCM_CREDS	send / receive descriptors send / receive user credentials

# Ancillary Data

- File descriptors and process credentials can be passed between unrelated processes using ancillary data.



# How much data is Queued?



- Use *recv()* with MSG\_PEEK flag.

```
2 int numbytes = recv(fd, buf, bufsize, MSG_PEEK);
```

- For TCP, this will give the number of bytes available in socket recv buffer.
  - This value could change in between two reads.
- For a connected UDP socket, this return the number of bytes in the next available datagram.
  - Between two reads this value remains same.
- Use *ioctl()* call with FIONREAD command.

```
2 ioctl(fd, FIONREAD, &numbytes)
```

- In UDP case, the size of datagram can be zero. This makes it difficult to distinguish between nodata or data.
  - Safer to use *select()* first and then call I/O.

# Sockets and Standard I/O




- TCP and UDP sockets are full duplex. File streams can also be full duplex.
- We can open a file stream on a socket using *fdopen()*.
- If we open with mode `r+`, then
  - An input function can't be followed by output function without calling `fseek()`.
  - An output function can't be followed by input function without calling `fseek()`.
  - But we can't call `Iseek()` on sockets.
- So open two separate streams on a socket: one for reading, one for writing.



# Unix Domain Sockets

# Internet Domain vs Unix Domain

- Internet Domain: *AF\_INET* or *AF\_INET6*
  - Used in network communication.
  - Can be used between two processes in the same host also.
- Unix Domain: *AF\_UNIX* or *AF\_LOCAL*
  - Used for communication between processes on the same host. Same API as sockets API.
  - No TCP/IP protocol stack. A socket is made of two buffers in the kernel.
  - No header processing, no checksums. Reliable communication. Unix domain sockets are twice faster.



# Unix Domain Sockets - Usage



- Unix domain sockets are used for three reasons:
  - Unix domain sockets are often twice as fast as a TCP socket when both peers are on the same host.
    - X Windows
  - used when passing file descriptors between processes on the same host.
  - unix domain sockets provide the client's process credentials (user ID and group IDs) to the server, which can provide additional security checking

# Unix Domain Sockets



- Two types of sockets are provided.
  - Stream sockets
    - Similar to TCP
  - Datagram Sockets
    - Similar to UDP sockets.
    - Message boundaries are preserved.
    - Communication is reliable unlike UDP.

# Unix Domain Socket Address



- End Point Address
  - pathnames within the normal file system
  - The pathname associated with a Unix domain socket should be an absolute pathname.

```
1 struct sockaddr_un {  
2     sa_family_t sun_family; /* Always AF_UNIX */  
3     char sun_path[108]; /* Null-terminated socket pathname */  
4 };
```

# Binding End Point to a Socket



- When used to bind a UNIX domain socket, bind() creates an entry in the file system.

```
1 const char *SOCKNAME = "/tmp/mysock";
2 int sfd;
3 struct sockaddr_un addr;
4 sfd = socket(AF_UNIX, SOCK_STREAM, 0); /* Create socket */
5 if (sfd == -1)
6     errExit("socket");
7 memset(&addr, 0, sizeof(struct sockaddr_un)); /* Clear structure */
8 addr.sun_family = AF_UNIX; /* UNIX domain address */
9 strncpy(addr.sun_path, SOCKNAME, sizeof(addr.sun_path) - 1);
10 if (bind(sfd, (struct sockaddr *) &addr, sizeof(struct sockaddr_un)) == -1)
11     errExit("bind");
```

- We can't bind a socket to an existing pathname (bind() fails with the error EADDRINUSE ).
- A socket may be bound to only one pathname; conversely, a pathname can be bound to only one socket.
  - When the socket is no longer required, its pathname entry should be removed using unlink().

# Unix Domain Stream Sockets



- Server
  - Absolute pathname is required.
  - Pathname specified in connect() should be existing, and bound to a socket
  - If the listening socket's queue is full, ECONREFUSED is immediately returned.
- Client
  - Create socket
  - Connect to the server.

# Unix Domain Stream Server

```
1  /*sockets/us_xfr_sv.c*/
2  main(int argc, char *argv[])
3  {   struct sockaddr_un addr;
4      int sfd, cfd;
5      ssize_t numRead;
6      char buf[BUF_SIZE];
7      sfd = socket(AF_UNIX, SOCK_STREAM, 0);
8      if (sfd == -1) errExit("socket");
9      if (remove(SV_SOCKET_PATH) == -1 && errno != ENOENT)
10         errExit("remove-%s", SV_SOCKET_PATH);
11     memset(&addr, 0, sizeof(struct sockaddr_un));
12     addr.sun_family = AF_UNIX;
13     strncpy(addr.sun_path, SV_SOCKET_PATH, sizeof(addr.sun_path) - 1);
14     if (bind(sfd, (struct sockaddr *) &addr, sizeof(struct sockaddr_un)) == -1)
15         errExit("bind");
16     if (listen(sfd, BACKLOG) == -1) errExit("listen");
17     for (;;) {
18         cfd = accept(sfd, NULL, NULL);
19         if (cfid == -1) errExit("accept");
20         while ((numRead = read(cfd, buf, BUF_SIZE)) > 0)
21             if (write(STDOUT_FILENO, buf, numRead) != numRead)
22                 fatal("partial/failed write");
23             if (numRead == -1) errExit("read");
24             if (close(cfd) == -1) errMsg("close");
25     }}
```

# Unix Domain Stream Client

```
1  /*sockets/us_xfr_cl.c*/
2  main(int argc, char *argv[])
3  {   struct sockaddr_un addr;
4      int sfd;
5      ssize_t numRead;
6      char buf[BUF_SIZE];
7      sfd = socket(AF_UNIX, SOCK_STREAM, 0); /* Create client socket */
8      if (sfd == -1) errExit("socket");
9      /* Construct server address, and make the connection */
10     memset(&addr, 0, sizeof(struct sockaddr_un));
11     addr.sun_family = AF_UNIX;
12     strncpy(addr.sun_path, SV_SOCKET_PATH, sizeof(addr.sun_path) - 1);
13     if (connect(sfd, (struct sockaddr *) &addr,
14                 sizeof(struct sockaddr_un)) == -1)
15         errExit("connect");
16     /* Copy stdin to socket */
17     while ((numRead = read(STDIN_FILENO, buf, BUF_SIZE)) > 0)
18         if (write(sfd, buf, numRead) != numRead)
19             fatal("partial/failed write");
20     if (numRead == -1)
21         errExit("read");
22     exit(EXIT_SUCCESS); /* Closes our socket; server sees EOF */
23 }
```

# Unix Domain Datagram Sockets

- Datagram sockets are reliable unlike UDP sockets.
  - Datagrams are not lost.
  - Datagrams are delivered in order and without duplicates.
- Server
  - Creates a socket
  - binds to well-known path.
- Client
  - Creates a socket
  - **binds the socket to an address, so that the server can send its reply.**
  - The client address is made unique by including the client's process ID in the pathname.

# Unix Domain Datagram Server



```
1  /*sockets/ud_udcase_sv.c*/
2  main(int argc, char *argv[])
3  {    struct sockaddr_un svaddr, claddr;
4      sfd = socket(AF_UNIX, SOCK_DGRAM, 0);           /* Create server socket */
5      if (remove(SV_SOCK_PATH) == -1 && errno != ENOENT)
6          errExit("remove-%s", SV_SOCK_PATH);
7      memset(&svaddr, 0, sizeof(struct sockaddr_un));
8      svaddr.sun_family = AF_UNIX;
9      strncpy(svaddr.sun_path, SV_SOCK_PATH, sizeof(svaddr.sun_path) - 1);
10     if (bind(sfd, (struct sockaddr *) &svaddr, sizeof(struct sockaddr_un)) == -1)
11         errExit("bind");
12     for (;;) {
13         len = sizeof(struct sockaddr_un);
14         numBytes = recvfrom(sfd, buf, BUF_SIZE, 0,
15                             (struct sockaddr *) &claddr, &len);
16         if (numBytes == -1) errExit("recvfrom");
17         printf("Server received %ld bytes from %s\n",
18                (long) numBytes,
19                claddr.sun_path);
20         for (j = 0; j < numBytes; j++)
21             buf[j] = toupper((unsigned char) buf[j]);
22         if (sendto(sfd, buf, numBytes, 0, (struct sockaddr *) &claddr, len) !=
23             numBytes)
24             fatal("sendto");
}}
```

# Unix Domain Datagram Client



```
1  /* sockets/ud_udcase_cl.c*/
2  main(int argc, char *argv[])
3  {   struct sockaddr_un svaddr, claddr;
4      sfd = socket(AF_UNIX, SOCK_DGRAM, 0);
5      memset(&claddr, 0, sizeof(struct sockaddr_un));
6      claddr.sun_family = AF_UNIX;
7      snprintf(claddr.sun_path, sizeof(claddr.sun_path),
8                 "/tmp/ud_udcase_cl.%ld", (long) getpid());
9      if (bind(sfd, (struct sockaddr *) &claddr, sizeof(struct sockaddr_un)) == -1)
10          errExit("bind");
11      /* Construct address of server */
12      memset(&svaddr, 0, sizeof(struct sockaddr_un));
13      svaddr.sun_family = AF_UNIX;
14      strncpy(svaddr.sun_path, SV_SOCKET_PATH, sizeof(svaddr.sun_path) - 1);
15      /* Send messages to server; echo responses on stdout */
16      for (j = 1; j < argc; j++) {
17          msgLen = strlen(argv[j]);           /* May be longer than BUF_SIZE */
18          if (sendto(sfd, argv[j], msgLen, 0, (struct sockaddr *) &svaddr,
19                     sizeof(struct sockaddr_un)) != msgLen)
20              fatal("sendto");
21          numBytes = recvfrom(sfd, resp, BUF_SIZE, 0, NULL, NULL);
22          if (numBytes == -1)    errExit("recvfrom");
23          printf("Response %d: %.*s\n", j, (int) numBytes, resp); }
24      remove(claddr.sun_path);/* Remove client socket pathname */
25 }
```


# socketpair()

- creates a pair of sockets and connect them together.

```
1 #include <sys/socket.h>
2 int socketpair(int domain, int type ,int protocol ,int sockfd [2]);
3 //Returns 0 on success, or -1 on error
```

- Returns two socket fds.
- No path names bound for sockets. Not visible outside the process.
- Type SOCK\_STREAM creates the equivalent of a bidirectional pipe (also known as a stream pipe).
- Each socket can be used for both reading and writing.
- Just like in pipe, after creating calling socketpair(), fork() is called.
- Parent and child can communicate using sockets.

# socketpair()



# socketpair()



```
1 #include <sys/types.h>
2 #include <unistd.h>
3 #include <sys/socket.h>
4 main ()
5 {
6     int i;
7     int p[2];
8     pid_t ret;
9     socketpair(AF_UNIX, SOCK_STREAM, 0, p);
10    ret = fork ();
11    if (ret == 0)
12    {
13        close (1);
14        dup (p[1]);
15        close (p[0]);
16        execlp ("ls", "ls", "-l", (char *) 0);
17    }
18    if (ret > 0)
19    {
20        close (0);
21        dup (p[0]);
22        close (p[1]);
23        execlp ("wc", "wc", "-l", (char *) 0);
24    }
25 }
```

# Passing File Descriptors

- Unix system provide a way to pass any open descriptor from one process to any other process.(using `sendmsg()`)
- It allows one process (typically a server) to do the privileged execution
  - dialing a modem, negotiating locks for the file or deal with database
  - simply pass back to the calling process a descriptor that can be used with all the I/O functions.
- All the details involved in opening the file or device are hidden from the client.


# Passing File Descriptors

- Steps involved
  - Create a unix domain socket(stream or datagram)
  - one process opens a descriptor by calling any of the unix function that returns a descriptor
  - the sending process build a ***msghdr*** structure containing the descriptor to be passed
  - Sending process sends ancillary data using ***sendmsg()*** with SCM\_RIGHTS
  - the receiving process calls ***recvmsg()*** to receive the descriptor on the unix domain socket
- Passing a descriptor is not same as passing descriptor number
  - involves creating a new descriptor in the receiving process that refers to the same file table entry within the kernel.

# Ancillary Data



- Ancillary data can be sent and received using the `msg_control` and `msg_controllen` members of the `msghdr` structure.
  - Another term for ancillary data is control information.



```
1 struct cmsghdr {  
2     socklen_t cmsg_len; /* length in bytes, including this structure */  
3     int cmsg_level; /* originating protocol */  
4     int cmsg_type; /* protocol-specific type */  
5     /* followed by unsigned char cmsg_data[] */  
6 };
```


# Ancillary Data

- Ancillary data is domain specific.

<b>Protocol</b>	<b>cmsg_level</b>	<b>Cmsg_type</b>	<b>Description</b>
IPv4	IPPROTO_IP	IP_RECVDSTADD R IP_RECVIF	receive destination address with UDP datagram receive interface index with UDP datagram
IPv6	IPPROTO_IPV6	IPV6_DSTOPTS IPV6_HOPLIMIT IPV6_HOPOPTS IPV6_NEXTHOP IPV6_PKTINFO IPV6_RTHDR	specify / receive destination options specify / receive hop limit specify / receive hop-by-hop options specify next-hop address specify / receive packet information specify / receive routing header
Unix domain	SOL_SOCKET	SCM_RIGHTS SCM_CREDS	send / receive descriptors send / receive user credentials

# Ancillary Data

- File descriptors and process credentials can be passed between unrelated processes using ancillary data.



# Passing File Descriptors Example

- Protocol
  - If  $\text{buf}[1] < 0$  then there is error.
  - If  $\text{buf}[1] = 0$  then it is success.
- Receiver
  - Create unix domain stream socket
  - Send file name
  - Call `recv_fd`
- Sender
  - Create unix domain stream socket
  - Open the file descriptor
  - Call `send_fd`

- Macros associated with ancillary data

```
1 #include <sys/socket.h>
2 #include <sys/param.h> /* for ALIGN macro on many implementations */
3 struct cmsghdr *CMSG_FIRSTHDR(struct msghdr *mhdrptr) ;
4 //Returns: pointer to first cmsghdr structure or NULL if no ancillary data
5 struct cmsghdr *CMSG_NXTHDR(struct msghdr *mhdrptr, struct cmsghdr *cmsgptr) ;
6 //Returns: pointer to next cmsghdr structure or NULL if no more ancillary data
7 unsigned char *CMSG_DATA(struct cmsghdr *cmsgptr) ;
8 //Returns: pointer to first byte of data associated with cmsghdr structure
9 unsigned int CMSG_LEN(unsigned int length) ;
10 //Returns: value to store in cmsg_len given the amount of data
11 unsigned int CMSG_SPACE(unsigned int length) ;
12 //Returns: total size of an ancillary data object given the amount of data
```

# Send fd

```
1 #include <sys/socket.h>
2 #define CONTROLLEN CMSG_LEN(sizeof(int))
3 static struct cmsghdr *cmptr = NULL;
4 int send_fd(int fd, int fd_to_send)
5 {   struct iovec iov[1];
6     struct msghdr msg;
7     char buf[2]; /* send_fd()/recv_fd() 2-byte protocol */
8     iov[0].iov_base = buf; iov[0].iov_len = 2;
9     msg.msg_iov = iov; msg.msg_iovlen = 1;
10    msg.msg_name = NULL; msg.msg_namelen = 0;
11    if (fd_to_send < 0) {
12        msg.msg_control = NULL;
13        msg.msg_controllen = 0;
14        buf[1] = -fd_to_send; /* nonzero status means error */
15    }
16    else { if (cmptr == NULL && (cmptr = malloc(CONTROLLEN)) == NULL)
17            return(-1);
18            cmptr->cmsg_level = SOL_SOCKET;
19            cmptr->cmsg_type = SCM_RIGHTS;
20            cmptr->cmsg_len = CONTROLLEN;
21            msg.msg_control = cmptr;
22            msg.msg_controllen = CONTROLLEN;
23            *(int *)CMSG_DATA(cmptr) = fd_to_send; /* the fd to pass */
24            buf[1] = 0; /* zero status means OK */
25        }
26        buf[0] = 0; /* null byte flag to recv_fd() */
27        if (sendmsg(fd, &msg, 0) != 2)
28            return(-1);
29        return(0);
30    }
```

# Receive fd



```
1 int recv_fd(int sockfd )
2 {
3     #define CONTROLLEN CMSG_LEN(sizeof(int))
4     static struct cmsghdr *cmptr = NULL;
5
6     struct iovec iov[1];
7     struct msghdr msg;
8     char buf[2]; /* send_fd()/recv_fd() 2-byte protocol */
9     memset(&msg, 0, sizeof(msg));
10    iov.iov_base = buf;
11    iov.iov_len = sizeof(data)-1;
12    msg.msg_iov = &iov;
13    msg.msg_iovlen = 1;
14    if (cmptr == NULL && (cmptr = malloc(CONTROLLEN)) == NULL)
15        return(-1);
16    msg.msg_control = cmptr;
17    msg.msg_controllen = CONTROLLEN;
18    recvmsg(sockfd, &msg, 0)
19    if (buf[1]<0) {
20        printf("failed to open %s: %s\n", name, data);
21        return -1;
22    }
23    /* Loop over all control messages */
24    cmsg = CMSG_FIRSTHDR(&msg);
25    while (cmsg != NULL) {
26        if (cmsg->cmsg_level == SOL_SOCKET
27            && cmsg->cmsg_type == SCM_RIGHTS)
28            return *(int *) CMSG_DATA(cmsg);
29        cmsg = CMSG_NXTHDR(&msg, cmsg);
30    }
31 }
```

# Passing Credentials

- A process can pass its credentials as ancillary data using SCM\_CREDS option.
- The structure for credentials

```
1 struct ucred {  
2     pid_t pid;      /* process ID of the sending process */  
3     uid_t uid;      /* user ID of the sending process */  
4     gid_t gid;      /* group ID of the sending process */  
5 };
```

- This structure is filled by the kernel and passed onto the receiver process.
- Example:
  - Sender process sends file name and access mode.
  - Server verifies the credentials and passes on the fd.

```
1 sockfd = socket(AF_LOCAL, SOCK_STREAM, 0);
2 bzero(&servaddr, sizeof(servaddr));
3 servaddr.sun_family = AF_LOCAL;
4 strcpy(servaddr.sun_path, "PATH");
5 connect(sockfd, (struct sockaddr *) &servaddr, sizeof(servaddr));
6 msgh.msg iov = &iov;
7 msgh.msg iovlen = 1;
8 /* Send Filename and Access Mode to server */
9 strcat(data, argv[1]);strcat(data, "#");
10 strcat(data, argv[2]);strcat(data, "#");
11 iov.iov_base = data;
12 iov.iov_len = MAX_DATA;
13 msgh.msg_name = NULL;
14 msgh.msg_namelen = 0;
15 msgh.msg_control = NULL;
16 msgh.msg_controllen = 0;
17
18 if(sendmsg(sockfd, &msgh, 0) < 0)
19 {
20     perror("Error sending message");
21     exit(1);
22 }
```

# Receiver



```
1 optval = 1;
2 /* Set SO_PASSCRED socket option for receiving credentials of other processes */
3 setsockopt(*(int *)arg, SOL_SOCKET, SO_PASSCRED, &optval, sizeof(optval));
4 /* Set 'control_un' to describe ancillary data that we want to receive */
5 control_un.cmh.cmsg_len = CMSG_LEN(sizeof(struct ucred));
6 control_un.cmh.cmsg_level = SOL_SOCKET;
7 control_un.cmh.cmsg_type = SCM_CREDENTIALS;
8 /* Set 'msgh' fields to describe 'control_un' */
9 msgh.msg_control = control_un.control;
10 msgh.msg_controllen = sizeof(control_un.control);
11 msgh.msg iov = &iov; msgh.msg iovlen = 1;
12 iov.iov_base = data;
13 iov.iov_len = MAX_DATA;
14 msgh.msg_name = NULL;
15 msgh.msg_namelen = 0;
16 /* Receive real plus ancillary data */
17 nr = recvmsg(*(int *)arg, &msgh, 0);
18 /* Extract credentials information from received ancillary data */
19 cmhp = CMSG_FIRSTHDR(&msgh);
20 ucredp = (struct ucred *) CMSG_DATA(cmhp);
21 printf("Received Credentials pid: %ld, uid: %ld, gid: %ld\n",
22 (long) ucredp->pid, (long) ucredp->uid, (long) ucredp->gid);
```

# The Linux Abstract Socket Namespace



- Is a Linux-specific feature that allows us to bind a UNIX domain socket to a name without that name being created in the file system.
- It is not necessary to unlink the socket pathname when we have finished using the socket. The abstract name is automatically removed when the socket is closed.
  - To create an abstract binding, we specify the first byte of the sun\_path field as a null byte ( \0 ).
  - The remaining bytes of the sun\_path field then define the abstract name for the socket. These bytes are interpreted in their entirety, rather than as a null-terminated string.

# The Linux Abstract Socket Namespace Example

```
1 struct sockaddr_un addr;
2 memset(&addr, 0, sizeof(struct sockaddr_un)); /* Clear address structure */
3 addr.sun_family = AF_UNIX; /* UNIX domain address */
4 /* addr.sun_path[0] has already been set to 0 by memset() */
5 strncpy(&addr.sun_path[1], "xyz", sizeof(addr.sun_path) - 2);
6 /* Abstract name is "xyz" followed by null bytes */
7 sockfd = socket(AF_UNIX, SOCK_STREAM, 0);
8 if (sockfd == -1)
9     errExit("socket");
10 if (bind(sockfd, (struct sockaddr *) &addr,
11           sizeof(struct sockaddr_un)) == -1)
12     errExit("bind");
```



# Unix I/O Models

# IO Models

- While doing I/O there are two phases
  - Waiting for the data
  - Copying the data
- Each I/O model differs how it deals with these two phases.
- There are five I/O models
  - blocking I/O
  - nonblocking I/O
  - I/O multiplexing (select and poll)
  - signal driven I/O (SIGIO)
  - asynchronous I/O (the POSIX aio\_functions)




# Blocking I/O Model

- Most prevalent model


# Nonblocking I/O Model

- When the socket is set to be non-blocking,
  - We tell the kernel that do not put the process to sleep if IO can't be completed.



# I/O Multiplexing Model

- Block in `select()` or `poll()` instead of blocking in actual I/O system call.




# I/O Multiplexing Model



- Blocking and I/O multiplexing seem to be non-different and actually calling two sys calls in I/O multiplexing.
- Advantage with I/O multiplexing is that it can wait for I/O on multiple fds.

# Signal-Driven I/O Model

- Tell the kernel to notify us with the SIGIO signal when the descriptor is ready.




# Signal-Driven I/O Model



- To use signal-driven I/O with a socket (SIGIO) requires the process to perform the following three steps:
  - A signal handler must be established for the SIGIO signal.
  - The socket owner must be set, normally with the F\_SETOWN command of fcntl.
  - Signal-driven I/O must be enabled for the socket, normally with the F\_SETFL command of fcntl to turn on the O\_ASYNC flag.

# Asynchronous I/O Model

- The main difference between this model and the signal-driven I/O models that
  - with signal-driven I/O, the kernel tells us when an I/O operation can be initiated,
  - but with asynchronous I/O, the kernel tells us when an I/O operation is complete.



# Asynchronous I/O Model


- POSIX API for asynchronous IO is implemented in a very few systems.
- *aiocb* structure

```
1 struct aiocb{  
2     int             aio_fildes //    file descriptor  
3     off_t          aio_offset //    file offset  
4     volatile void* aio_buf   //    location of buffer  
5     size_t         aio_nbytes //    length of transfer  
6     int            aio_reqprio //    request priority offset  
7     struct sigevent aio_sigevent // signal number and value  
8     int            aio_lio_opcode //operation to be performed  
9 };
```

```
1 #include <aio.h>  
2 int aio_read(struct aiocb *aiocbp);
```

# Comparison of I/O Models

- First four models differ only in first phase.



# Synchronous I/O vs Asynchronous I/O

- POSIX defines these two terms as follows:
  - A synchronous I/O operation causes the requesting process to be blocked until that I/O operation completes.
  - An asynchronous I/O operation does not cause the requesting process to be blocked.
- Using these definitions,
  - the first four I/O models
    - blocking,
    - nonblocking,
    - I/O multiplexing,
    - and signal-driven I/O
  - are all synchronous because the actual I/O operation (`recvfrom`) blocks the process.
  - Only the asynchronous I/O model matches the asynchronous I/O definition.



# I/O Multiplexing

T1:Ch6

# I/O Multiplexing

- I/O multiplexing allows us to simultaneously monitor multiple file descriptors to see if I/O is possible on any of them.
- `select()`, appeared along with the sockets API in BSD. This was historically the more widespread of the two system calls. The other system call, `poll()`, appeared in System V.
- We can use `select()` and `poll()` to monitor file descriptors for regular files, terminals, pseudoterminals, pipes, FIFOs, sockets, and some types of character devices.
- Both system calls allow a process either to block indefinitely waiting for file descriptors to become ready or to specify a timeout on the call.

# select()



- The select() system call blocks until one or more of a set of file descriptors becomes ready.

```
1 #include <sys/time.h> /* For portability */
2 #include <sys/select.h>
3 int select(int nfds , fd_set * readfds , fd_set * writefds,
4           fd_set * exceptfds, struct timeval * timeout );
5 //Returns number of ready file descriptors, 0 on timeout, or -1 on error
```

- nfds*: highest number assigned to a descriptor +1.
- readfds*: set of descriptors we want to read from.
- writefds*: set of descriptors we want to write to.
- exceptfds*: set of descriptors to watch for exceptions.
- timeout*: maximum time select should wait

```
7 struct timeval {
8     long tv_sec;          /* seconds */
9     long tv_usec;         /* microseconds */
10 }
```

# select()

- **timeval==NULL**
  - Wait forever : return only when descriptor is ready
- **timeval != NULL:** wait up to a fixed amount of time
  - timeval = 0
    - Do not wait at all : return immediately after checking the descriptors
  - Timeval>0
    - Return only if descriptor is ready or timeval expires.

# File descriptor sets

- The readfds, writefds, and exceptfds arguments are pointers to file descriptor sets, represented using the data type *fd\_set*.
- the *fd\_set* data type is implemented as a bit mask.

```
1 #include <sys/select.h>
2 void FD_ZERO(fd_set *fdset);
3 /* clear all bits in fdset */
4 void FD_SET(int fd, fd_set *fdset);
5 /* turn on the bit for fd in fdset */
6 void FD_CLR(int fd, fd_set *fdset);
7 /* turn off the bit for fd in fdset */
8 int FD_ISSET(int fd, fd_set *fdset);
9 /* is the bit for fd on in fdset ? */
10 //Returns true (1) if fd is in fdset, or false (0) otherwise
```

- A file descriptor set has a maximum size, defined by the constant `FD_SETSIZE`. On Linux, this constant has the value 1024.

# select()

- *nfds*
  - Its value is the maximum descriptor to be tested, plus one
    - example: fds 1,2,5 => nfds: 6
- *readset*
  - descriptor set for checking readable
- *writeset*
  - descriptor set for checking writable
- *exceptset*
  - descriptor set for checking two exception conditions
    - arrival of out of band data for a socket
    - the presence of control status information to be read from the master side of a pseudo terminal
- When select returns value>1, these sets have been modified by kernel. Now they contain the fds which are ready.

# When is the descriptor ready for reading?


- The number of bytes of data in the socket receive buffer is greater than or equal to the current size of the low-water mark for the socket receive buffer. SO\_RCVLOWAT socket option. It defaults to 1 for TCP and UDP sockets
- The read half of the connection is closed (i.e., a TCP connection that has received a FIN)
- The socket is a listening socket and the number of completed connections is nonzero.
- A socket error is pending. A read operation on the socket will not block and will return an error (-1) with errno set to the specific error condition.

# When the socket is ready for writing?

- The number of bytes of available space in the socket send buffer is greater than or equal to the current size of the low-water mark for the socket send buffer. 2048 bytes.
- The write half of the connection is closed. A write operation on the socket will generate SIGPIPE.
- A socket using a non-blocking connect has completed the connection, or the connect has failed
- A socket error is pending. A write operation on the socket will not block and will return an error (-1) with errno set to the specific error condition.
- These pending errors can also be fetched and cleared by calling getsockopt with the SO\_ERROR socket option.

# Client Handling Multiple Descriptors

- A client is handling two descriptors.
  - *stdin*
  - *socket*
- Sequential handling:
  - First wait on *stdin*.
  - Write to *socket*
  - Read from *socket*.
  - Write to *stdout*.




Normal course of actions

```


1 void str_cli(FILE *fp, int sockfd)
2 {
3     char sendline[MAXLINE], recvline[MAXLINE];
4     while (Fgets(sendline, MAXLINE, fp) != NULL) {
5         Writen(sockfd, sendline, strlen (sendline));
6         if (Readline(sockfd, recvline, MAXLINE) == 0)
7             err_quit("str_cli: server terminated prematurely");
8         Fputs(recvline, stdout);
9     }
10 }
```

# Client Handling Multiple Descriptors

- `read()` call on both `stdin` and socket will block until data is available.
- Consider a case:
  - If client is blocked in waiting for user to enter data, meanwhile TCP receives `FIN` from server.
    - Server is down. So sending request is meaningless.
- How to handle uncertainty of availability of data on descriptors?




Normal course of actions



Unexpected Server shutdown

# read() on socket

- Peer TCP sends data, the socket becomes readable and *read* returns greater than 0.
- Peer TCP send a FIN(peer process terminates), the socket become readable and *read* returns 0(end-of-file)
- Peer TCP send a RST(peer host has crashed and rebooted), the socket become readable and returns -1 and *errno* contains the specific error code



# Client Handling Multiple Descriptors

- To avoid a situation where data has arrived from socket but client is unable to take note of it, use I/O Multiplexing.
- Client can wait on select().
  - Add stdin, socket to fd\_set.
  - Call select () with fd\_set for readability
  - When select() returns, find out which descriptor is ready with data.
  - Call read() on that fd.
- This will enable client to give timely response and avoid error situations.


# Client Handling Multiple Descriptors

```
1 void str_cli(FILE *fp, int sockfd)
2 {
3     int maxfdp1;
4     fd_set rset;
5     char sendline[MAXLINE], recvline[MAXLINE];
6     FD_ZERO(&rset);
7     for ( ; ; ) {
8         FD_SET(fileno(fp), &rset);
9         FD_SET(sockfd, &rset);
10        maxfdp1 = max(fileno(fp), sockfd) + 1;
11        select(maxfdp1, &rset, NULL, NULL, NULL);
12        if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
13            if (Readline(sockfd, recvline, MAXLINE) == 0)
14                err_quit("str_cli: server terminated prematurely");
15            Fputs(recvline, stdout);
16        }
17        if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
18            if (Fgets(sendline, MAXLINE, fp) == NULL)
19                return; /* all done */
20            Writen(sockfd, sendline, strlen(sendline));
21        }
22    }//for
23 } //str_cli
```



# Batch Mode Client

- Stop and Wait client
  - Send one request and wait for reply.
  - Usual in interactive mode.
- Batch Mode clients
  - Send requests without waiting for reply.
  - Better bandwidth utilization.

Time 7:



Time 8:



# Batch Mode Client

- Need for closing a socket partially:
  - We tell server that we have sent all requests by closing socket.
    - It will send FIN to server.
  - But in batch mode, by closing socket, we send FIN but we can't read replies which are yet to reach the client.
- `close()` vs `shutdown()` sys calls
  - closes the socket partially (either read end or write end) unlike `close()` sys call.
    - `close()` closes completely.
  - Irrespective of reference count it closes the socket.
    - `close()` will initiate FIN only if reference count for the fd reaches 0.

# shutdown() sys call

- Sometimes, it is useful to close one half of the connection, so that data can be transmitted in just one direction through the socket.

```
1 #include <sys/socket.h>
2 int shutdown(int sockfd , int how );
3 //Returns 0 on success, or -1 on error
```

- SHUT\_RD : read-half of the connection closed. Subsequent reads will return end-of-file (0).
  - SHUT\_RD can't be used meaningfully for TCP sockets.
- SHUT\_WR : write-half of the connection closed. Also called *socket half-close*. Buffered data will be sent followed by termination sequence.
  - Common use of shutdown()
  - Subsequent writes to the local socket yield the SIGPIPE signal and an EPIPE error.
- SHUT\_RDWR : both closed
  - Note that shutdown() doesn't close the file descriptor, even if how is specified as SHUT\_RDWR . To close the file descriptor, we must additionally call close().


# Batch Mode Client



- After user presses, Ctrl-D (EOF), close write half of the socket.
- Also set `stdineof` variable to 1.
  - This will help in inferring the FIN received from server as normal or abnormal termination.
  - In case of normal termination, we received all the replies.


EOF on `stdin`  
Set `stdineof=1`  
`shutdown(fd, SHUT_WR)`

*if stdineof=1  
Receiving FIN  
from server is  
Normal  
termination*



Data on `stdin`

*if stdineof=0  
Receiving FIN  
from server is  
abnormal  
termination*



```
1 str_cli(FILE *fp, int sockfd)
2 {
3     int      maxfdp1, stdineof;
4     fd_set  rset;
5     stdineof = 0;
6     FD_ZERO(&rset);
7     for ( ; ; ) {
8         if (stdineof == 0)
9             FD_SET(fileno(fp), &rset);
10    FD_SET(sockfd, &rset);
11    maxfdp1 = max(fileno(fp), sockfd) + 1;
12    select(maxfdp1, &rset, NULL, NULL, NULL);
13    if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
14        if ( (n = Read(sockfd, buf, MAXLINE)) == 0) {
15            if (stdineof == 1)
16                return;          /* normal termination */
17            else
18                err_quit("str_cli: server terminated prematurely");
19        }
20        Write(fileno(stdout), buf, n);
21    }
22    if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
23        if ( (n = Read(fileno(fp), buf, MAXLINE)) == 0) {
24            stdineof = 1;
25            shutdown(sockfd, SHUT_WR); /* send FIN */
26            FD_CLR(fileno(fp), &rset);
27            continue;
28        }
29        Writen(sockfd, buf, n);
30    }
31 }
```

# TCP Server Using select()




- Single process server that uses select to handle any number of clients, instead of forking one child per client.
- Protocol: echo
- Two data structures:
  - Client array
    - Keeps list of client sockets connected currently
  - fd\_set *allset*
    - Keeps list of fds for checking against readability.

# TCP Server Using select()

- There are three fds: 0,1,2
- One more fd after creating listening socket.


Before first client has established a connection



# TCP Server Using select()

- When a new client is accepted through accept()
  - A connected socket is added

After first client connection is established




# TCP Server Using select()



- When second client is accepted


After second client connection is established



# TCP Server Using select()

- When the first client terminates connection
  - This is known when read() returns zero.

After first client terminates its connection



- \*Maxfd does not change
- \* fd3 => listening socket fd
- \* fd4 => client1 socket fd deleted
- \* fd5 => client2 socket fd

# TCP Server Using select()

- Create a passive socket.

```
1 int main(int argc, char **argv)
2 {
3     int          i, maxi, maxfd, listenfd, connfd, sockfd;
4     int          nready, client[FD_SETSIZE];
5     fd_set       rset, allset;
6     struct sockaddr_in cliaddr, servaddr;
7     listenfd = socket(AF_INET, SOCK_STREAM, 0);
8     bzero(&servaddr, sizeof(servaddr));
9     servaddr.sin_family      = AF_INET;
10    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
11    servaddr.sin_port        = htons(SERV_PORT);
12    bind(listenfd, (SA *) &servaddr, sizeof(servaddr));
13    listen(listenfd, LISTENQ);
```

# TCP Server Using select()



- Handling when listening socket is readable

```
1  maxfd = listenfd;           /* initialize */
2  maxi = -1;                 /* index into client[] array */
3  for (i = 0; i < FD_SETSIZE; i++)
4    client[i] = -1;           /* -1 indicates available entry */
5  FD_ZERO(&allset);
6  FD_SET(listenfd, &allset);
7  for ( ; ; ) {
8    rset = allset;           /* structure assignment */
9    nready = select(maxfd+1, &rset, NULL, NULL, NULL);

10 if (FD_ISSET(listenfd, &rset)) /* new client connection */
11   clilen = sizeof(cliaddr);
12   connfd = accept(listenfd, (SA *) &cliaddr, &clilen);
13   for (i = 0; i < FD_SETSIZE; i++)
14     if (client[i] < 0) {
15       client[i] = connfd;    /* save descriptor */
16       break;
17     if (i == FD_SETSIZE) err_quit("too many clients");
18     FD_SET(connfd, &allset); /* add new descriptor to set */
19     if (connfd > maxfd)
20       maxfd = connfd;       /* maxfd for select */
21     if (i > maxi)
22       maxi = i;             /* max index in client[] array */
23     if (--nready <= 0)
24       continue;             /* no more readable descriptors */
25 }
```

# TCP Server Using select()

- When a connected socket is readable

```
1  for (i = 0; i <= maxi; i++) /* check all clients for data */
2      if ((sockfd = client[i]) < 0)
3          continue;
4      if (FD_ISSET(sockfd, &rset)) {
5          if ( (n = Readline(sockfd, line, MAXLINE)) == 0) {
6              /*connection closed by client */
7              close(sockfd);
8              FD_CLR(sockfd, &allset);
9              client[i] = -1;
10         }
11     else
12         Writen(sockfd, line, n);
13     if (--nready <= 0)
14         break; /* no more readable descriptors */
15   }
16 }
```

- This code looks complicated when compared to fork-per-client model. But this design avoids overhead of fork().

# Denial-of-Service Attacks



- If malicious client connect to the server, send 1 byte of data (other than a newline), and then goes to sleep.
  - in readline(), server is blocked.
- Solution
  - use nonblocking I/O
  - have each client serviced by a separate thread of control (spawn a process or a thread to service each client)
  - place a timeout on the I/O operation

# pselect()



```
1 struct timespec{  
2     time_t tv_sec; /*seconds*/  
3     long   tv_nsec; /* nanoseconds */  
4 };
```

- pselect contains two changes from the normal select function:
  - pselect uses the timespec structure instead of the timeval structure.
  - Accepts signal mask.

```
1 #define _XOPEN_SOURCE 600  
2 #include <sys/select.h>  
3 int pselect(int nfds , fd_set * readfds , fd_set * writefds ,  
4             fd_set * exceptfds,struct timespec * timeout , const sigset_t * sigmask );  
5 //Returns number of ready file descriptors, 0 on timeout, or -1 on error
```

```
2 ready = pselect(nfds, &readfds, &writefds, &exceptfds, timeout, &sigmask);
```

- This call is equivalent to

```
1 sigset_t origmask;  
2 sigprocmask(SIG_SETMASK, &sigmask, &origmask);  
3 ready = select(nfds, &readfds, &writefds, &exceptfds, timeout);  
4 sigprocmask(SIG_SETMASK, &origmask, NULL); /* Restore signal mask */
```

# Problems with select() and poll()

- The select() and poll() system calls are the portable, long-standing, and widely used methods of monitoring multiple file descriptors for readiness.
- Suffer from some problems
  - Kernel must check all the fds to check if they are ready.
  - Each time select() passes data structures which kernel modifies and returns.
  - Once select() returns, the program must inspect the data structure to see which fds are ready.
- Select() scales poorly with the increase of *fds*.
- Signal driven I/O or *epoll* (event poll) provide a scalable solution.



# Non-blocking I/O on Sockets

# Socket Operations

- Input operations: read, recv, ready, recvfrom, recvmsg
  - Blocking operations
  - TCP: until a byte arrives.
  - UDP: until a datagram arrives.
  - With non-blocking socket, if no data, return with EWOULDBLOCK error.
- Output Operations: write, send, writev, sendto, sendmsg
  - Blocks if there is no room in socket send buffer.
  - TCP: until all the data is written.
  - UDP: no send buffer present.
  - With non-blocking socket, TCP write will write whatever it can and returns no. of bytes written. If no room at all, it returns with error EWOULDBLOCK.


# Socket Operations



- Accepting incoming connections: `accept`
  - Blocks if no incoming connection.
  - With non-blocking socket, it would return with an error.
- Initiating Connections: `connect`
  - Blocks until client TCP receives ACK.
  - With non-blocking socket, it returns `errno EINPROGRESS`, and continues to establish connection.

# Client Handling a socket, stdin, stdout

- A client usually deals with
  - Stdin
  - Stdout
  - Socket



# Client Handling a socket, stdin, stdout

- We looked at
  - Stop and wait client
  - Batch mode client - select with blocking I/O
    - Once select() returns, and if socket is readable, read() is called on socket.
    - readline() call gets blocked on socket till it gets required data.
    - During this time, other clients have to wait.
- Now we look at select with non-blocking I/O
  - In this, read() will be a non-blocking operation. It will read whatever data available on socket. It returns.
  - When this fd is readable next time, further data is read.
    - This requires that we track the number of bytes read and the pointers in the buffer.

```
1 str_cli(FILE *fp, int sockfd)
2 {
3     int      maxfdp1, stdineof;
4     fd_set  rset;
5     stdineof = 0;
6     FD_ZERO(&rset);
7     for ( ; ; ) {
8         if (stdineof == 0)
9             FD_SET(fileno(fp), &rset);
10    FD_SET(sockfd, &rset);
11    maxfdp1 = max(fileno(fp), sockfd) + 1;
12    select(maxfdp1, &rset, NULL, NULL, NULL);
13    if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
14        if ( (n = Read(sockfd, buf, MAXLINE)) == 0) {
15            if (stdineof == 1)
16                return;          /* normal termination */
17            else
18                err_quit("str_cli: server terminated prematurely");
19        }
20        Write(fileno(stdout), buf, n);
21    }
22    if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
23        if ( (n = Read(fileno(fp), buf, MAXLINE)) == 0) {
24            stdineof = 1;
25            shutdown(sockfd, SHUT_WR); /* send FIN */
26            FD_CLR(fileno(fp), &rset);
27            continue;
28        }
29        Writen(sockfd, buf, n);
30    }
31 }
```

# Select with Non-Blocking IO



- Non-blocking IO complicates buffer management.
  - We have to keep track of how much is read and how much is written.
- Two buffers:
  - *to*: reading from standard input and write to socket.
  - *from*: read from socket and write to stdout.




Figure 15.1 Buffer containing data from standard input going to the socket.

# Select with Non-Blocking IO




Figure 15.2 Buffer containing data from the socket going to standard output.

- **foptr**: points to the next byte to be sent to `stdout`.
- **friptr**: points to the next byte into which next byte can be read..

# Select with Non-Blocking IO



```
1 void str_cli(FILE *fp, int sockfd)
2 {
3     int      maxfdp1, val, stdineof;
4     ssize_t n, nwritten;
5     fd_set rset, wset;
6     char    to[MAXLINE], fr[MAXLINE];
7     char    *toiptr, *tooptr, *friptr, *froptr;
8     val = fcntl(sockfd, F_GETFL, 0);
9     fcntl(sockfd, F_SETFL, val | O_NONBLOCK);
10    val = fcntl(STDIN_FILENO, F_GETFL, 0);
11    fcntl(STDIN_FILENO, F_SETFL, val | O_NONBLOCK);
12    val = fcntl(STDOUT_FILENO, F_GETFL, 0);
13    fcntl(STDOUT_FILENO, F_SETFL, val | O_NONBLOCK);
14    toiptr = tooptr = to;           /* initialize buffer pointers */
15    friptr = froptr = fr;
16    stdineof = 0;
17
18    maxfdp1 = max(max(STDIN_FILENO, STDOUT_FILENO), sockfd) + 1;
19    for ( ; ; ) {
20        FD_ZERO(&rset);
21        FD_ZERO(&wset);
22        if (stdineof == 0 && toiptr < &to[MAXLINE])
23            FD_SET(STDIN_FILENO, &rset);      /* read from stdin */
24        if (friptr < &fr[MAXLINE])
25            FD_SET(sockfd, &rset); /* read from socket */
26        if (tooptr != toiptr)
27            FD_SET(sockfd, &wset); /* data to write to socket */
28        if (froptr != friptr)
29            FD_SET(STDOUT_FILENO, &wset); /* data to write to stdout */
30        select(maxfdp1, &rset, &wset, NULL, NULL);
```

# reads from standard input

```

30     select(maxfdp1, &rset, &wset, NULL, NULL);
31     if (FD_ISSET(STDIN_FILENO, &rset)) {
32       if((n = read(STDIN_FILENO, toiptr, &to[MAXLINE] - toiptr)) < 0) {
33         if (errno != EWOULDBLOCK)
34           err_sys("read error on stdin");
35       } else if (n == 0) {
36         fprintf(stderr, "%s: EOF on stdin\n", gf_time());
37         stdineof = 1; /* all done with stdin */
38         if (toiptr == toiptr)
39           shutdown(sockfd, SHUT_WR); /* send FIN */
40       } else {
41         fprintf(stderr, "%s: read %d bytes from stdin\n", gf_time(),
42                 n);
43         toiptr += n; /* # just read */
44         FD_SET(sockfd, &wset); /* try and write to socket below */
45     }
46   }
  
```

Amt of space available in to buffer



If user has pressed Ctrl-D,  
set stdineof=1  
If no outstanding data on  
buffer, close the write end.

Increment to pointer  
set socket in wset for writability

# reads from socket

```

47 if (FD_ISSET(sockfd, &rset)) {
48   if ((n = read(sockfd, friptr, &fr[MAXLINE] - friptr)) < 0) {
49     if (errno != EWOULDBLOCK)
50       err_sys("read error on socket");
51   } else if (n == 0) {
52     fprintf(stderr, "%s: EOF on socket\n", gf_time());
53     if (stdineof)
54       return; /* normal termination */
55     else
56       err_quit("str_cli: server terminated prematurely");
57   } else {
58     fprintf(stderr, "%s: read %d bytes from socket\n",
59             gf_time(), n);
60     friptr += n; /* # just read */
61     FD_SET(STDOUT_FILENO, &wset); /* try and write below */
62   }
63 }
```

Amt of space available in from buffer

Increment friptr.  
Add stdout to wset to test for writability.

# writes to standard output

```
65 if (FD_ISSET(STDOUT_FILENO, &wset) && ((n = friptr - froptr) > 0)) {  
66     if ( (nwritten = write(STDOUT_FILENO, froptr, n)) < 0) {  
67         if (errno != EWOULDBLOCK)  
68             err_sys("write error to stdout");  
69     } else {  
70         fprintf(stderr, "%s: wrote %d bytes to stdout\n",  
71                 gf_time(), nwritten);  
72         froptr += nwritten; /* # just written */  
73         if (froptr == friptr)  
74             froptr = friptr = fr; /* back to beginning of buffer */  
75     }  
76 }
```

No of bytes to write >0

If the write is successful, froptr is incremented by the number of bytes written

# Writes to socket

No of bytes to write >0

```
78 if (FD_ISSET(sockfd, &wset) && ((n = toiptr - tooptr) > 0)) {  
79     if ( (nwritten = write(sockfd, tooptr, n)) < 0) {  
80         if (errno != EWOULDBLOCK)  
81             err_sys("write error to socket");  
82     } else {  
83         fprintf(stderr, "%s: wrote %d bytes to socket\n",  
84                 gf_time(), nwritten);  
85         tooptr += nwritten; /* # just written */  
86         if (tooptr == toiptr) {  
87             toiptr = tooptr = to; /* back to beginning of buffer */  
88             if (stdineof)  
89                 shutdown(sockfd, SHUT_WR); /* send FIN */  
90         }  
91     }  
92 }  
93 }  
94 }
```

If the write is successful, tooptr is incremented by the number of bytes written


if we encountered an EOF on standard input, the FIN can be sent to the server



# Client using Multiple Processes

# Client using Multiple Processes

- Whenever we find the need to use nonblocking I/O, it will usually be simpler to split the application into either processes (using fork) or threads.
  - Parent reads from stdin and writes to socket
  - Child reads from socket and writes to stdout.



- Normal termination occurs when the EOF on standard input is encountered. The parent reads this EOF and calls shutdown to send a FIN.
- If abnormal occurs, the child will read an EOF on the socket. If this happens, the child must tell the parent to stop copying from the standard input to the socket
  - the child sends a signal (e.g. SIGTERM) to the parent.

# Client using Multiple Processes

```
1 void str_cli(FILE *fp, int sockfd)
2 {
3     pid_t      pid;
4     char      sendline[MAXLINE], recvline[MAXLINE];
5     if ( (pid = fork()) == 0) { /* child: server -> stdout */
6         while (Readline(sockfd, recvline, MAXLINE) > 0)
7             fputs(recvline, stdout);
8         kill(getppid(), SIGTERM); /* in case parent still running */
9         exit(0);
10    }
11    /* parent: stdin -> server */
12    while (fgets(sendline, MAXLINE, fp) != NULL)
13        Writen(sockfd, sendline, strlen(sendline));
14    shutdown(sockfd, SHUT_WR); /* EOF on stdin, send FIN */
15    pause();
16    return;
17 }
```

# Comparing Client Designs

- when copying 2,000 lines from a client to a server with an RTT of 175 ms:

Client	Time taken for sending and receiving
stop-and-wait	354.0 sec
select and blocking I/O	12.3 sec
nonblocking I/O	6.9 sec
fork	8.7 sec
threaded version	8.5 sec

- nonblocking I/O version is almost twice as fast as version using blocking I/O with select.
- Version using fork is slower than nonblocking I/O version.
- Nevertheless, given the complexity of the nonblocking I/O code versus the fork code, fork version is simple approach.



# Non-blocking Connect

# Nonblocking connect()

- TCP socket nonblocking connect
  - return: an error of EINPROGRESS
  - TCP three-way handshake continues
  - check the connection establishment using select
- There are three uses for a nonblocking connect.
  - We can overlap other processing with the three-way handshake.
  - We can establish multiple connections at the same time using this technique.
    - popular with Web browsers
  - Since we wait for the connection establishment to complete using select, we can specify time limit for select, allowing us to shorten the timeout for the connect.

# Nonblocking connect()

- Set the socket to non-blocking.
- Call `connect()`. It will return immediately with error `EINPROGRESS`.
- We use `select()` to check what has happened to `connect()`.
- If the descriptor is readable or writable, we call `getsockopt()` to fetch the socket's pending error (`SO_ERROR`). If the connection completed successfully, this value will be 0.

# Nonblocking connect()



```
1 int connect_nonb(int sockfd, const SA *saptr, socklen_t salen, int nsec)
2 {
3     int      flags, n, error;
4     socklen_t len;
5     fd_set rset, wset;
6     struct timeval tval;
7     flags = fcntl(sockfd, F_GETFL, 0);
8     fcntl(sockfd, F_SETFL, flags | O_NONBLOCK);
9     error = 0;
10    if ( (n = connect(sockfd, saptr, salen)) < 0)
11        if (errno != EINPROGRESS)
12            return (-1);
13    /* Do whatever we want while the connect is taking place. */
14    if (n == 0)
15        goto done;                  /* connect completed immediately */
16    FD_ZERO(&rset);
17    FD_SET(sockfd, &rset);
18    wset = rset;
19    tval.tv_sec = nsec;
20    tval.tv_usec = 0;
```

# Nonblocking connect()

```
22     if ( (n = select(sockfd + 1, &rset, &wset, NULL,
23             nsec ? &tval : NULL)) == 0) {
24         close(sockfd);           /* timeout */
25         errno = ETIMEDOUT;
26         return (-1);
27     }
28     if (FD_ISSET(sockfd, &rset) || FD_ISSET(sockfd, &wset)) {
29         len = sizeof(error);
30         if (getsockopt(sockfd, SOL_SOCKET, SO_ERROR, &error, &len) < 0)
31             return (-1);        /* Solaris pending error */
32     } else
33         err_quit("select error: sockfd not set");
34 done:
35     Fcntl(sockfd, F_SETFL, flags); /* restore file status flags */
36     if (error) {
37         close(sockfd);          /* just in case */
38         errno = error;
39         return (-1);
40     }
41     return (0);
42 }
```






# Web client: Non-blocking Connect

# nonblocking connect: web client



- A real-world example of nonblocking connects started with Netscape Web Client
- The client establishes an HTTP connection with a Web server and fetches a home page.
- On that page are often numerous references to other Web pages.
- Instead of fetching these other pages serially, one at a time, the client can fetch more than one at the same time, using nonblocking connects.

# nonblocking connect: web client (cont.)



# Non-blocking Connect

- This program will read up to 20 files from a Web server.
- We specify as command-line arguments
  - the maximum number of parallel connections,
  - the server's hostname, and
  - each of the filenames to fetch from the server.

```
2 bash$ web 3 www.foobar.com image1.gif image2.gif image3.gif image4.gif
3           image5.gif image6.gif image7.gif
```

- It means
  - three simultaneous connection
  - server's hostname
  - filename for the home page
  - the files to be read

# nonblock/web.h

```
1 #define MAXFILES      20
2 #define SERV          "80"           /* port number or service name */
3 struct file {
4     char   *f_name;             /* filename */
5     char   *f_host;            /* hostname or IPv4/IPv6 address */
6     int    f_fd;               /* descriptor */
7     int    f_flags;             /* F_xxx below */
8 } file[MAXFILES];
9 #define F_CONNECTING    1           /* connect() in progress */
10 #define F_READING       2           /* connect() complete; now reading */
11 #define F_DONE          4           /* all done */
12 #define GET_CMD         "GET %s HTTP/1.0\r\n\r\n"
13 /* globals */
14 int   nconn, nfiles, nlefttoconn, nlefttoread, maxfd;
15 fd_set rset, wset;
16 /* function prototypes */
17 void  home_page(const char *, const char *);
18 void  start_connect(struct file *);
19 void  write_get_cmd(struct file *);
```

- Each file has a state and fd.

```
1 main(int argc, char **argv)
2 {
3     int      i, fd, n, maxnconn, flags, error;
4     char    buf[MAXLINE];
5     fd_set  rs, ws;
6     if (argc < 5)
7         err_quit("usage: web <#conns> <hostname> <homepage> <file1> ...");
8     maxnconn = atoi(argv[1]);
9     nfiles = min(argc - 4, MAXFILES);
10    for (i = 0; i < nfiles; i++) {
11        file[i].f_name = argv[i + 4];
12        file[i].f_host = argv[2];
13        file[i].f_flags = 0;
14    }
15    printf("nfiles = %d\n", nfiles);
16    home_page(argv[2], argv[3]);
17    FD_ZERO(&rset);
18    FD_ZERO(&wset);
19    maxfd = -1;
20    nlefttoread = nlefttoconn = nfiles;
21    nconn = 0;
```

- Process command-line arguments
- Read home page
- Initialize globals
  - Fd sets, nconn is current number of connections.

# nonblock/start\_connect.c

```

1 void start_connect(struct file *fptr)
2 {
3     int      fd, flags, n;
4     ai = Host_serv(fptr->f_host, SERV, 0, SOCK_STREAM);
5     fd = Socket(ai->ai_family, ai->ai_socktype, ai->ai_protocol);
6     fptr->f_fd = fd;
7     printf("start_connect for %s, fd %d\n", fptr->f_name, fd);
8     /* Set socket nonblocking */
9     flags = Fcntl(fd, F_GETFL, 0);
10    Fcntl(fd, F_SETFL, flags | O_NONBLOCK);
11    /* Initiate nonblocking connect to the server. */
12    if ((n = connect(fd, ai->ai_addr, ai->ai_addrlen)) < 0) {
13        if (errno != EINPROGRESS)
14            err_sys("nonblocking connect error");
15        fptr->f_flags = F_CONNECTING;
16        FD_SET(fd, &rset);      /* select for reading and writing */
17        FD_SET(fd, &wset);
18        if (fd > maxfd)
19            maxfd = fd;
20    } else if (n >= 0)           /* connect is already done */
21        write_get_cmd(fptr);    /* write() the GET command */
22 }

```

- Initiate nonblocking connect
- Handle connection complete
- If connect returns successfully, the connection is already complete and the function write\_get\_cmd ends a command to the server.

# nonblock/write\_get\_cmd.c

```
1 #include "web.h"
2 void write_get_cmd(struct file *fptr)
3 {
4     int n;
5     char line[MAXLINE];
6     n = snprintf(line, sizeof(line), GET_CMD, fptr->f_name);
7     Writen(fptr->f_fd, line, n);
8     printf("wrote %d bytes for %s\n", n, fptr->f_name);
9     fptr->f_flags = F_READING; /* clears F_CONNECTING */
10    FD_SET(fptr->f_fd, &rset); /* will read server's reply */
11    if (fptr->f_fd > maxfd)
12        maxfd = fptr->f_fd;
13 }
```

- Build command and send it
- Set flags

# Main function: web.c

```
1 while (nlefttoread > 0) {
2   while (nconn < maxnconn && nlefttoconn > 0) {
3     /* find a file to read */
4     for (i = 0; i < nfiles; i++)
5       if (file[i].f_flags == 0)
6         break;
7     if (i == nfiles)
8       err_quit("nlefttoconn = %d but nothing found", nlefttoconn);
9     start_connect(&file[i]);
10    nconn++;
11    nlefttoconn--;
12 }
```

- Initiate another connection, if possible

# Main function: web.c

```

13     rs = rset;
14     ws = wset;
15     n = select(maxfd + 1, &rs, &ws, NULL, NULL);
16     for (i = 0; i < nfiles; i++) {
17         flags = file[i].f_flags;
18         if (flags == 0 || flags & F_DONE)
19             continue;
20         fd = file[i].f_fd;
21         if (flags & F_CONNECTING &&
22             (FD_ISSET(fd, &rs) || FD_ISSET(fd, &ws))) {
23             n = sizeof(error);
24             if (getsockopt(fd, SOL_SOCKET, SO_ERROR, &error, &n) < 0 ||
25                 error != 0) {
26                 err_ret("nonblocking connect failed for %s",
27                         file[i].f_name);
28             }
29             /* connection established */
30             printf("connection established for %s\n", file[i].f_name);
31             FD_CLR(fd, &wset); /* no more writability test */
32             write_get_cmd(&file[i]); /* write() the GET command */

```

- select waits for either readability or writability.
  - Descriptors that have a nonblocking connect in progress will be enabled in both sets, while descriptors with a completed connection that are waiting for data from the server will be enabled in just the read set.

# Main function: web.c

```
29     /* connection established */
30     printf("connection established for %s\n", file[i].f_name);
31     FD_CLR(fd, &wset); /* no more writeability test */
32     write_get_cmd(&file[i]); /* write() the GET command */
33 } else if (flags & F_READING && FD_ISSET(fd, &rs)) {
34     if ((n = Read(fd, buf, sizeof(buf))) == 0) {
35         printf("end-of-file on %s\n", file[i].f_name);
36         Close(fd);
37         file[i].f_flags = F_DONE; /* clears F_READING */
38         FD_CLR(fd, &rset);
39         nconn--;
40         nlefttoread--;
41     } else {
42         printf("read %d bytes from %s\n", n, file[i].f_name);
43     }
44 }
45 }
46 }
47 exit(0);
48 }
```

- If the F\_READING flag is set and the descriptor is ready for reading, we call read.

# Performance of Nonblocking Connect

- Table shows the clock time required to fetch a Web server's home page, followed by nine image files from that server.
  - The RTT to the server is about 150 ms.
  - The home page size was 4,017 bytes and the average size of the 9 image files was 1,621 bytes.
  - TCP's segment size was 512 bytes.
- Most of the improvement is obtained with three simultaneous connections.

# simultaneous connections	Clock time (seconds), non blocking	Clock time(seconds) Threads
1	6.0	6.3
2	4.1	4.2
3	3.0	3.1
4	2.8	3.0
5	2.5	2.7
6	2.4	2.5
7	2.3	2.3
8	2.2	2.3
9	2.0	2.2

# Q&A





# Thank You